# Managing infrastructure on EC2 Spot

Peter Juriz

# What we do

- Modelling Platform for various forms of data:
- Financial / Insurance / Logistical / Medical/ Marketing
- Compute, Simulation, Machine Learning and Data Visualization
- Modelling has a visual component

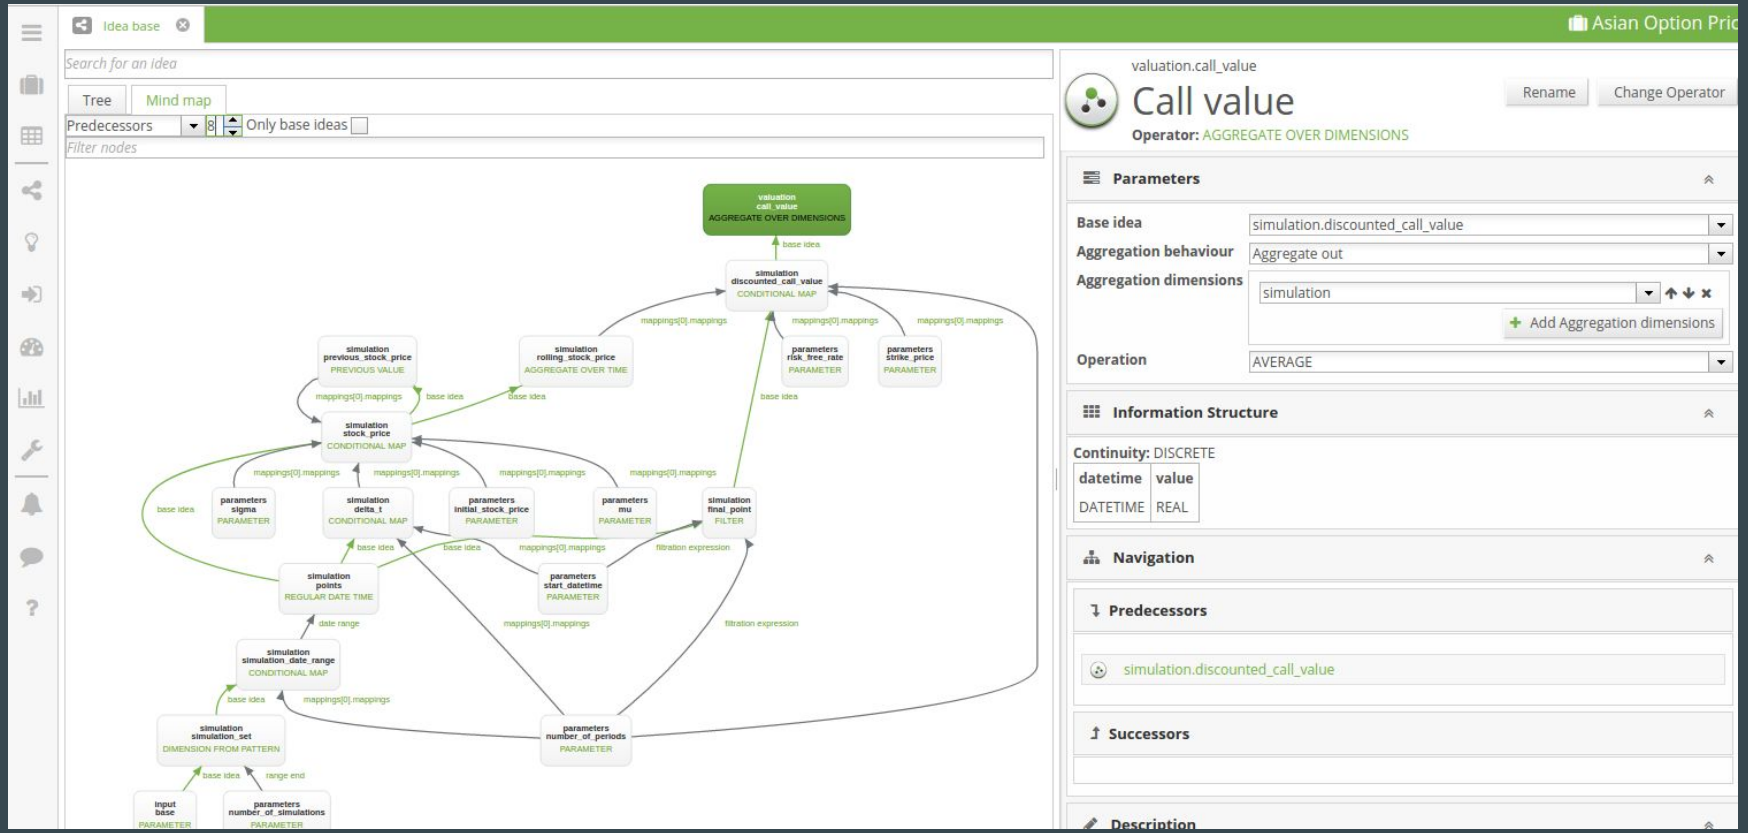


# Graph based compute engine (C++)

Compile these graphs into dependent Hadoop Map-Reduce jobs

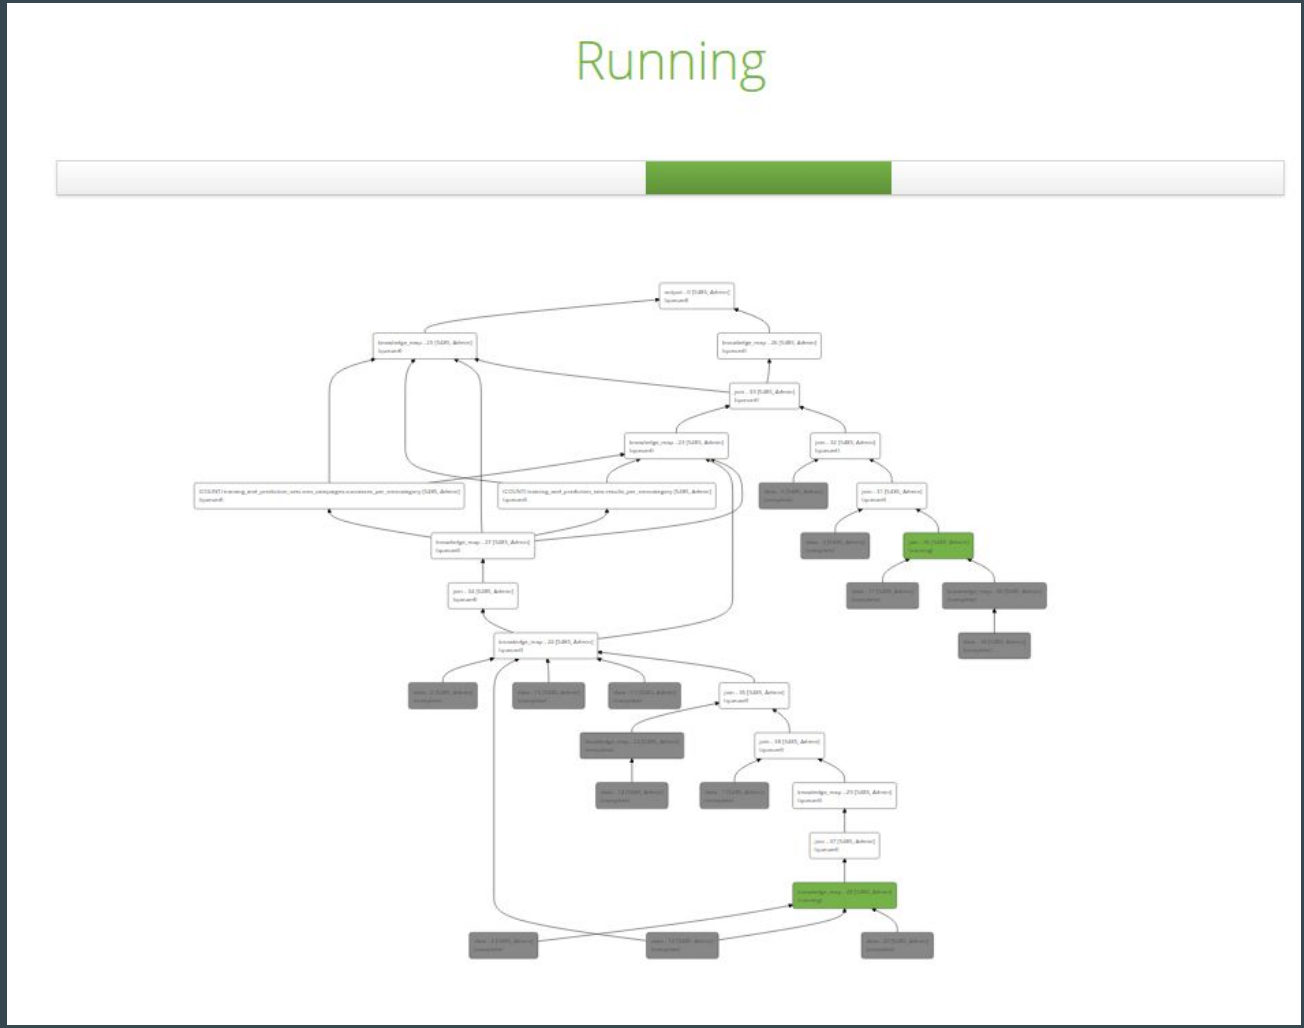

# Calculations can get pretty big


CSSH: hadoop05
CSSH: hadoop06
CSSH: hadoop07
CSSH: hadoop09
CSSH: hadoop10
CSSH: hadoop12
CSSH: hadoop13

# Motivations for using EC2 Spot

- Big calculations require many instances
- Much cheaper ~ $0.24 spot vs $1.06 on demand for r4.4xlarge per hour (4.4 times cheaper) - when not surging
- Can handle losing some machines

# How Spot Pricing Works

- Quasi market
- Set a bid, pay the market price
- If market price goes above your bid => your instance gets terminated
- Spot price can go 10X above the on demand price
- Different prices for different instance types

# Different volatility for different instance types

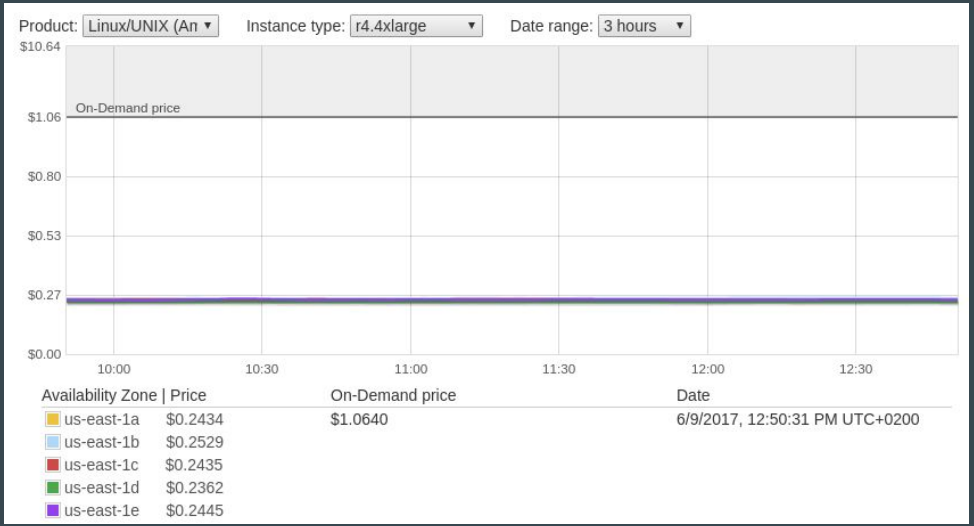


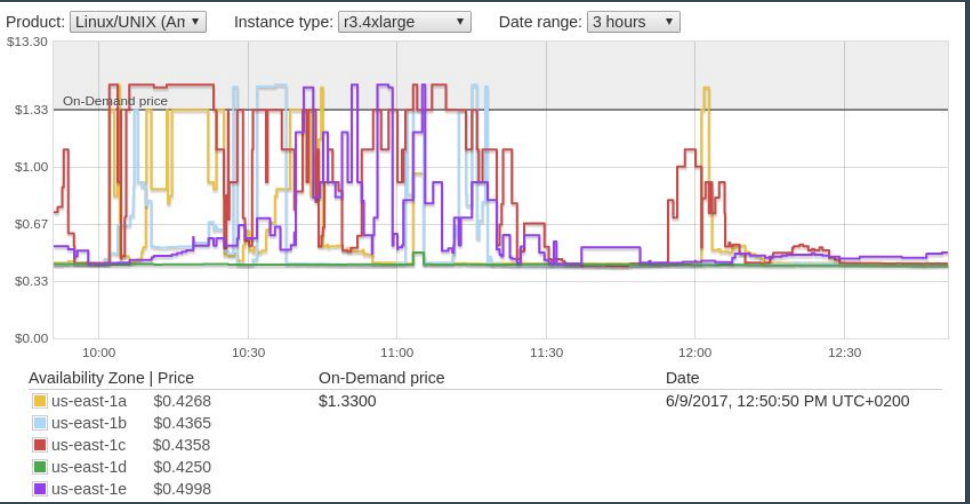

# Jumps in market price

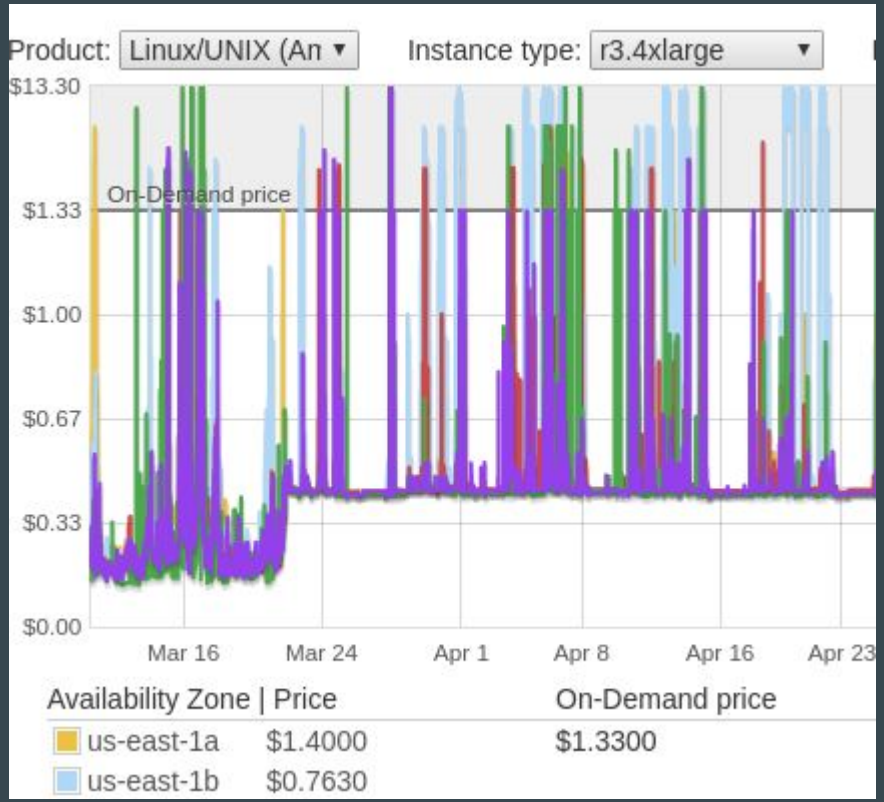

# Problems we've encountered

- What if price stays up?
- Multi-AZ - incurs interzone traffic costs
- Losing local disk can be a problem
- Some services don't handle node churn

# Having insights is important

# Build internal monitoring tools

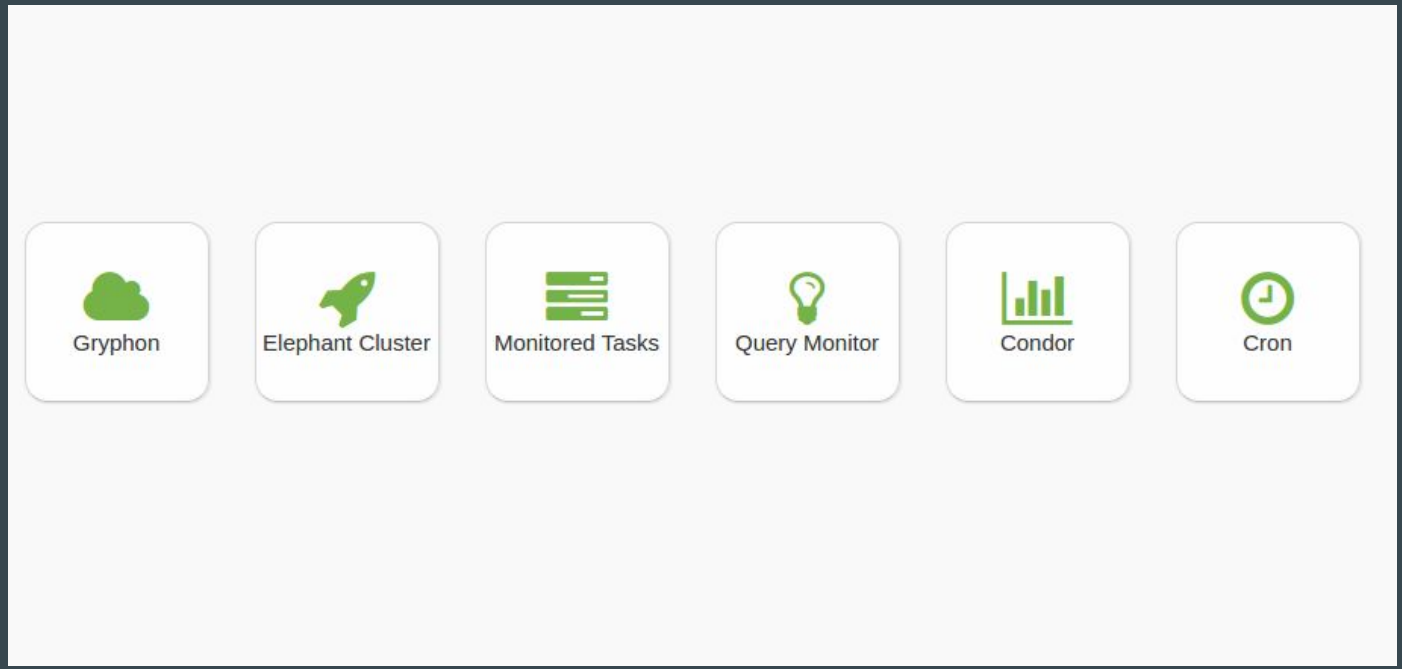

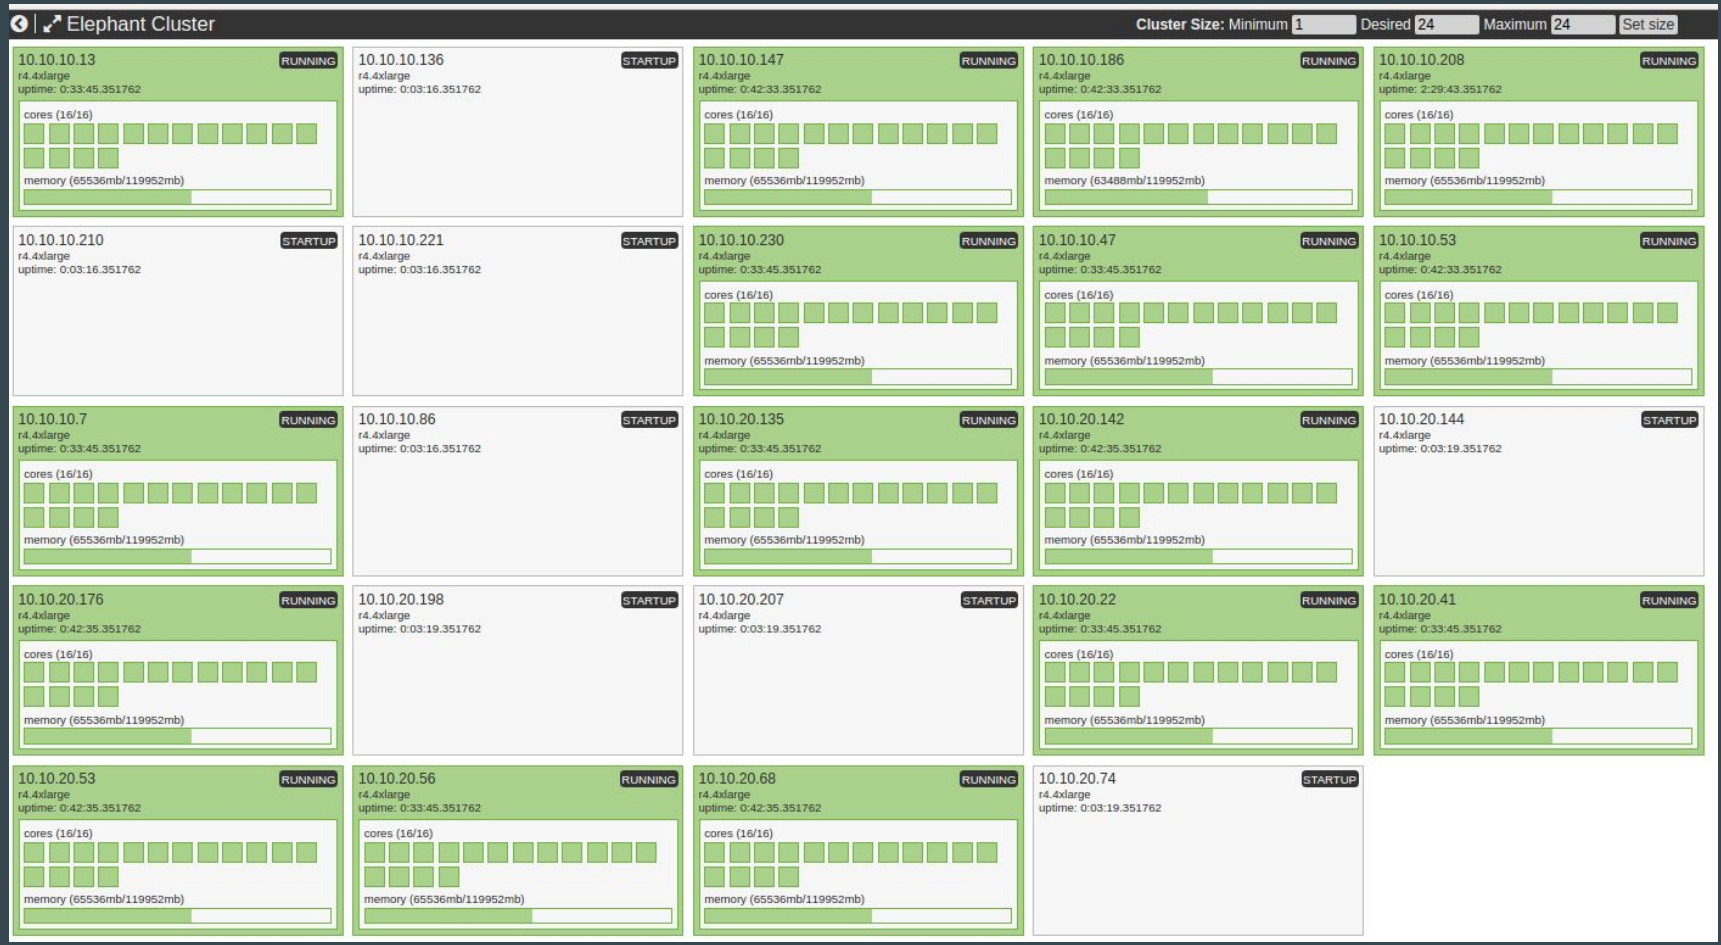
Elephant Cluster
Cluster Size: Minimum 1 Desired 24 Maximum 24 Set size
10.10.10.13 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.136 STARTUP
r4.4xlarge
uptime: 0:03:16.351762
10.10.10.147 RUNNING
r4.4xlarge
uptime: 0:42:33.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.186 RUNNING
r4.4xlarge
uptime: 0:42:33.351762
cores (16/16)
memory (63488mb/119952mb)
10.10.10.208 RUNNING
r4.4xlarge
uptime: 2:29:43.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.210 STARTUP
r4.4xlarge
uptime: 0:03:16.351762
10.10.10.221 STARTUP
r4.4xlarge
uptime: 0:03:16.351762
10.10.10.230 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.47 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.53 RUNNING
r4.4xlarge
uptime: 0:42:33.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.7 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.10.86 STARTUP
r4.4xlarge
uptime: 0:03:16.351762
10.10.20.135 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.142 RUNNING
r4.4xlarge
uptime: 0:42:35.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.144 STARTUP
r4.4xlarge
uptime: 0:03:19.351762
10.10.20.176 RUNNING
r4.4xlarge
uptime: 0:42:35.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.198 STARTUP
r4.4xlarge
uptime: 0:03:19.351762
10.10.20.207 STARTUP
r4.4xlarge
uptime: 0:03:19.351762
10.10.20.22 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.41 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.53 RUNNING
r4.4xlarge
uptime: 0:42:35.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.56 RUNNING
r4.4xlarge
uptime: 0:33:45.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.68 RUNNING
r4.4xlarge
uptime: 0:42:35.351762
cores (16/16)
memory (65536mb/119952mb)
10.10.20.74 STARTUP
r4.4xlarge
uptime: 0:03:19.351762

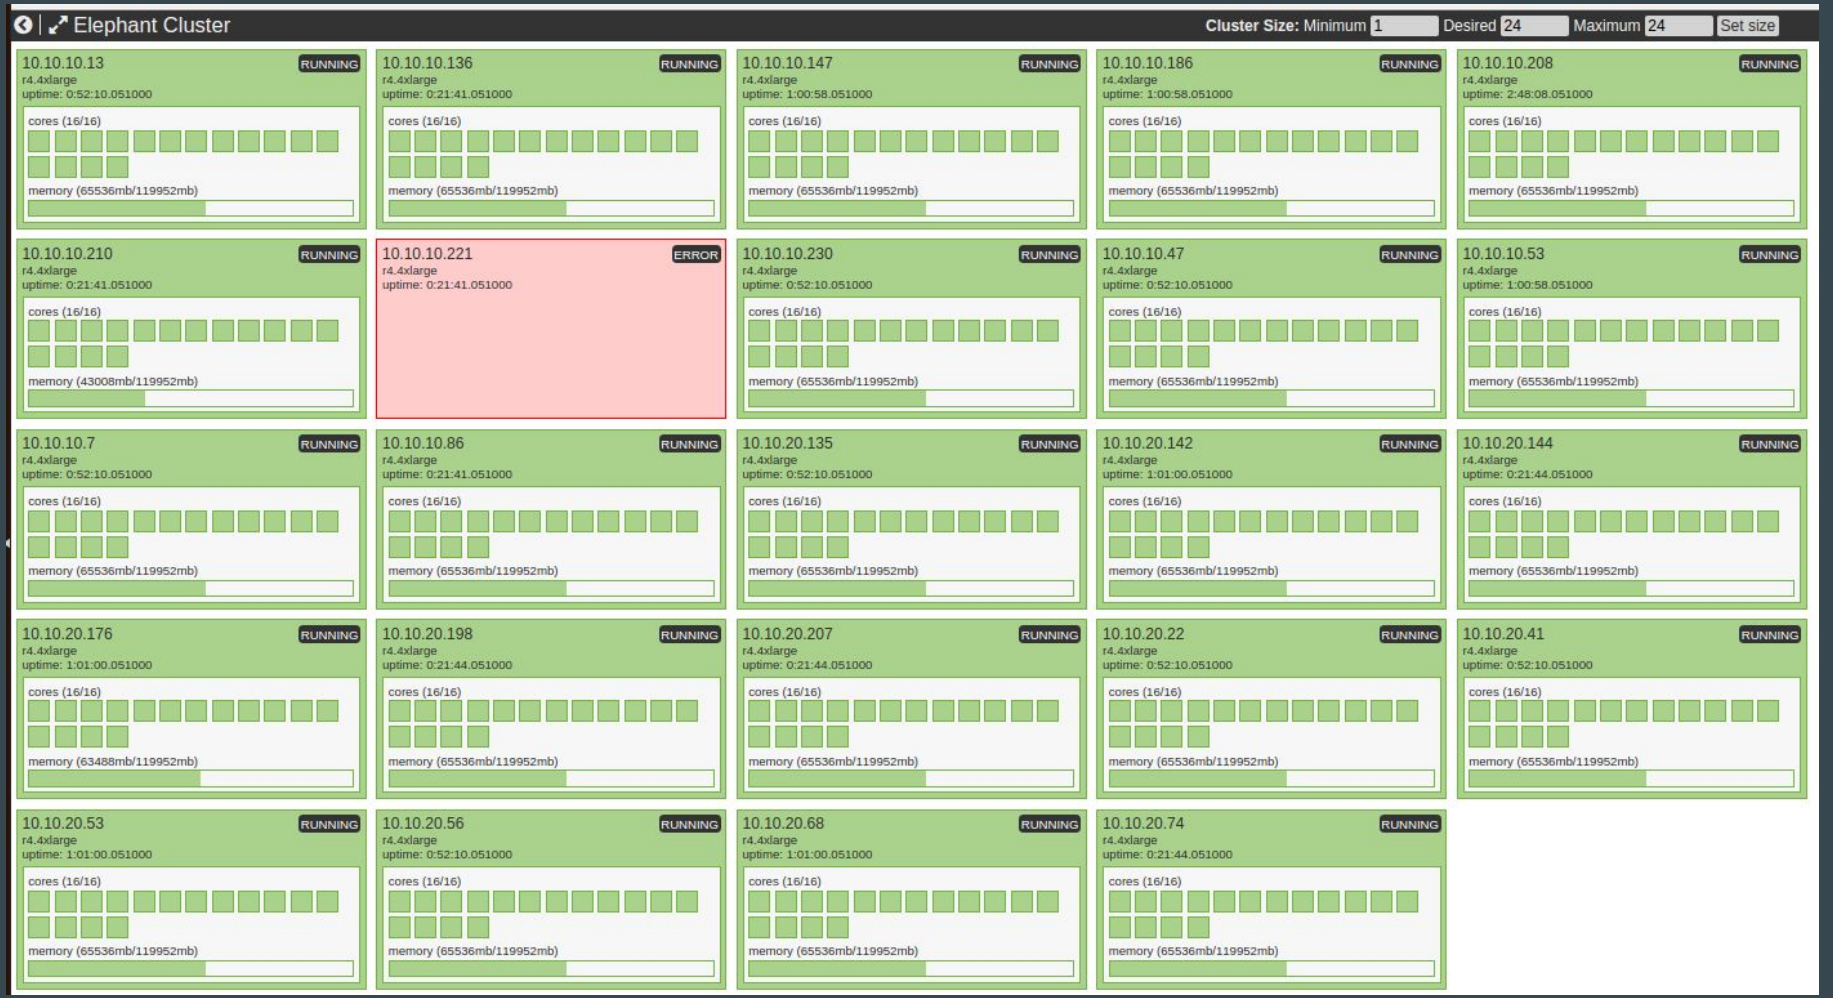
Elephant Cluster
Cluster Size: Minimum 1 Desired 24 Maximum 24 Set size
10.10.10.13 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.136 RUNNING
r4.4xlarge
uptime: 0:21:41.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.147 RUNNING
r4.4xlarge
uptime: 1:00:58.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.186 RUNNING
r4.4xlarge
uptime: 1:00:58.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.208 RUNNING
r4.4xlarge
uptime: 2:48:08.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.210 RUNNING
r4.4xlarge
uptime: 0:21:41.051000
cores (16/16)
memory (43008mb/119952mb)
10.10.10.221 ERROR
r4.4xlarge
uptime: 0:21:41.051000
10.10.10.230 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.47 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.53 RUNNING
r4.4xlarge
uptime: 1:00:58.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.7 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.10.86 RUNNING
r4.4xlarge
uptime: 0:21:41.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.135 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.142 RUNNING
r4.4xlarge
uptime: 1:01:00.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.144 RUNNING
r4.4xlarge
uptime: 0:21:44.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.176 RUNNING
r4.4xlarge
uptime: 1:01:00.051000
cores (16/16)
memory (63488mb/119952mb)
10.10.20.198 RUNNING
r4.4xlarge
uptime: 0:21:44.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.207 RUNNING
r4.4xlarge
uptime: 0:21:44.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.22 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.41 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.53 RUNNING
r4.4xlarge
uptime: 1:01:00.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.56 RUNNING
r4.4xlarge
uptime: 0:52:10.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.68 RUNNING
r4.4xlarge
uptime: 1:01:00.051000
cores (16/16)
memory (65536mb/119952mb)
10.10.20.74 RUNNING
r4.4xlarge
uptime: 0:21:44.051000
cores (16/16)
memory (65536mb/119952mb)

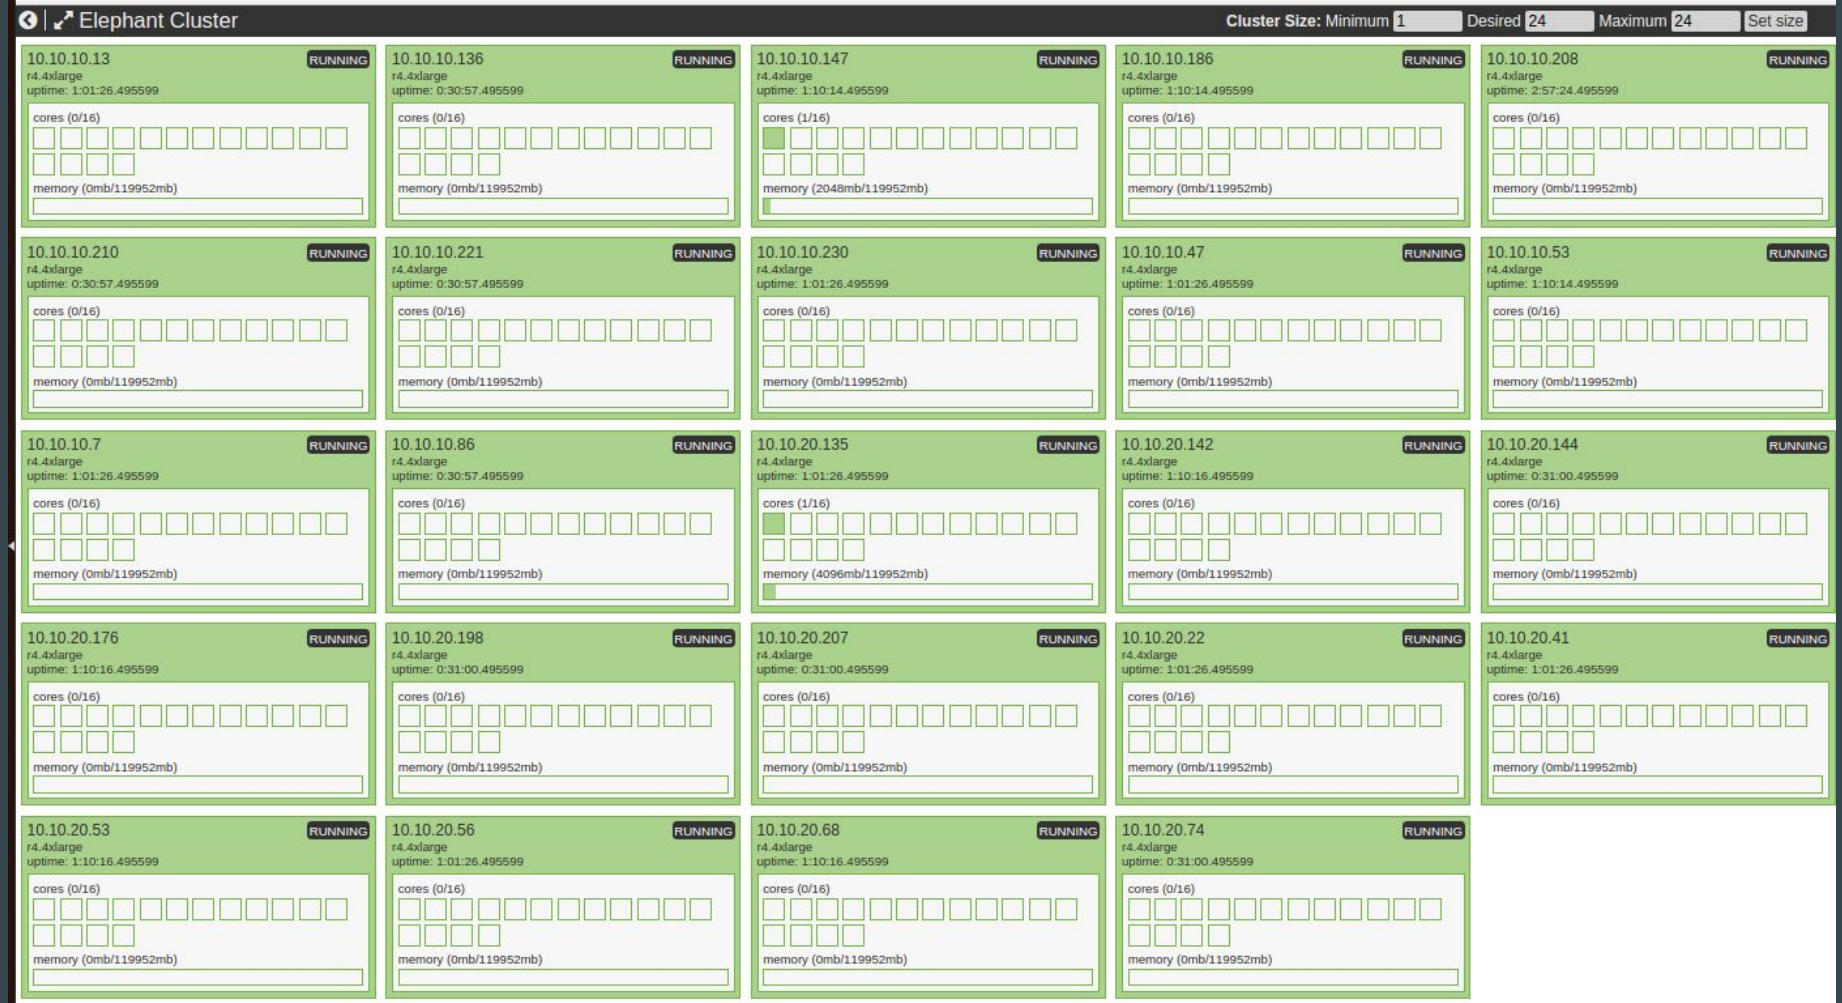
Elephant Cluster
Cluster Size: Minimum 1 Desired 24 Maximum 24 Set size
10.10.10.13 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.136 RUNNING
r4.4xlarge
uptime: 0:30:57.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.147 RUNNING
r4.4xlarge
uptime: 1:10:14.495599
cores (1/16)
memory (2048mb/119952mb)
10.10.10.186 RUNNING
r4.4xlarge
uptime: 1:10:14.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.208 RUNNING
r4.4xlarge
uptime: 2:57:24.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.210 RUNNING
r4.4xlarge
uptime: 0:30:57.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.221 RUNNING
r4.4xlarge
uptime: 0:30:57.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.230 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.47 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.53 RUNNING
r4.4xlarge
uptime: 1:10:14.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.7 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.10.86 RUNNING
r4.4xlarge
uptime: 0:30:57.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.135 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (1/16)
memory (4096mb/119952mb)
10.10.20.142 RUNNING
r4.4xlarge
uptime: 1:10:16.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.144 RUNNING
r4.4xlarge
uptime: 0:31:00.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.176 RUNNING
r4.4xlarge
uptime: 1:10:16.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.198 RUNNING
r4.4xlarge
uptime: 0:31:00.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.207 RUNNING
r4.4xlarge
uptime: 0:31:00.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.22 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.41 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.53 RUNNING
r4.4xlarge
uptime: 1:10:16.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.56 RUNNING
r4.4xlarge
uptime: 1:01:26.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.68 RUNNING
r4.4xlarge
uptime: 1:10:16.495599
cores (0/16)
memory (0mb/119952mb)
10.10.20.74 RUNNING
r4.4xlarge
uptime: 0:31:00.495599
cores (0/16)
memory (0mb/119952mb)

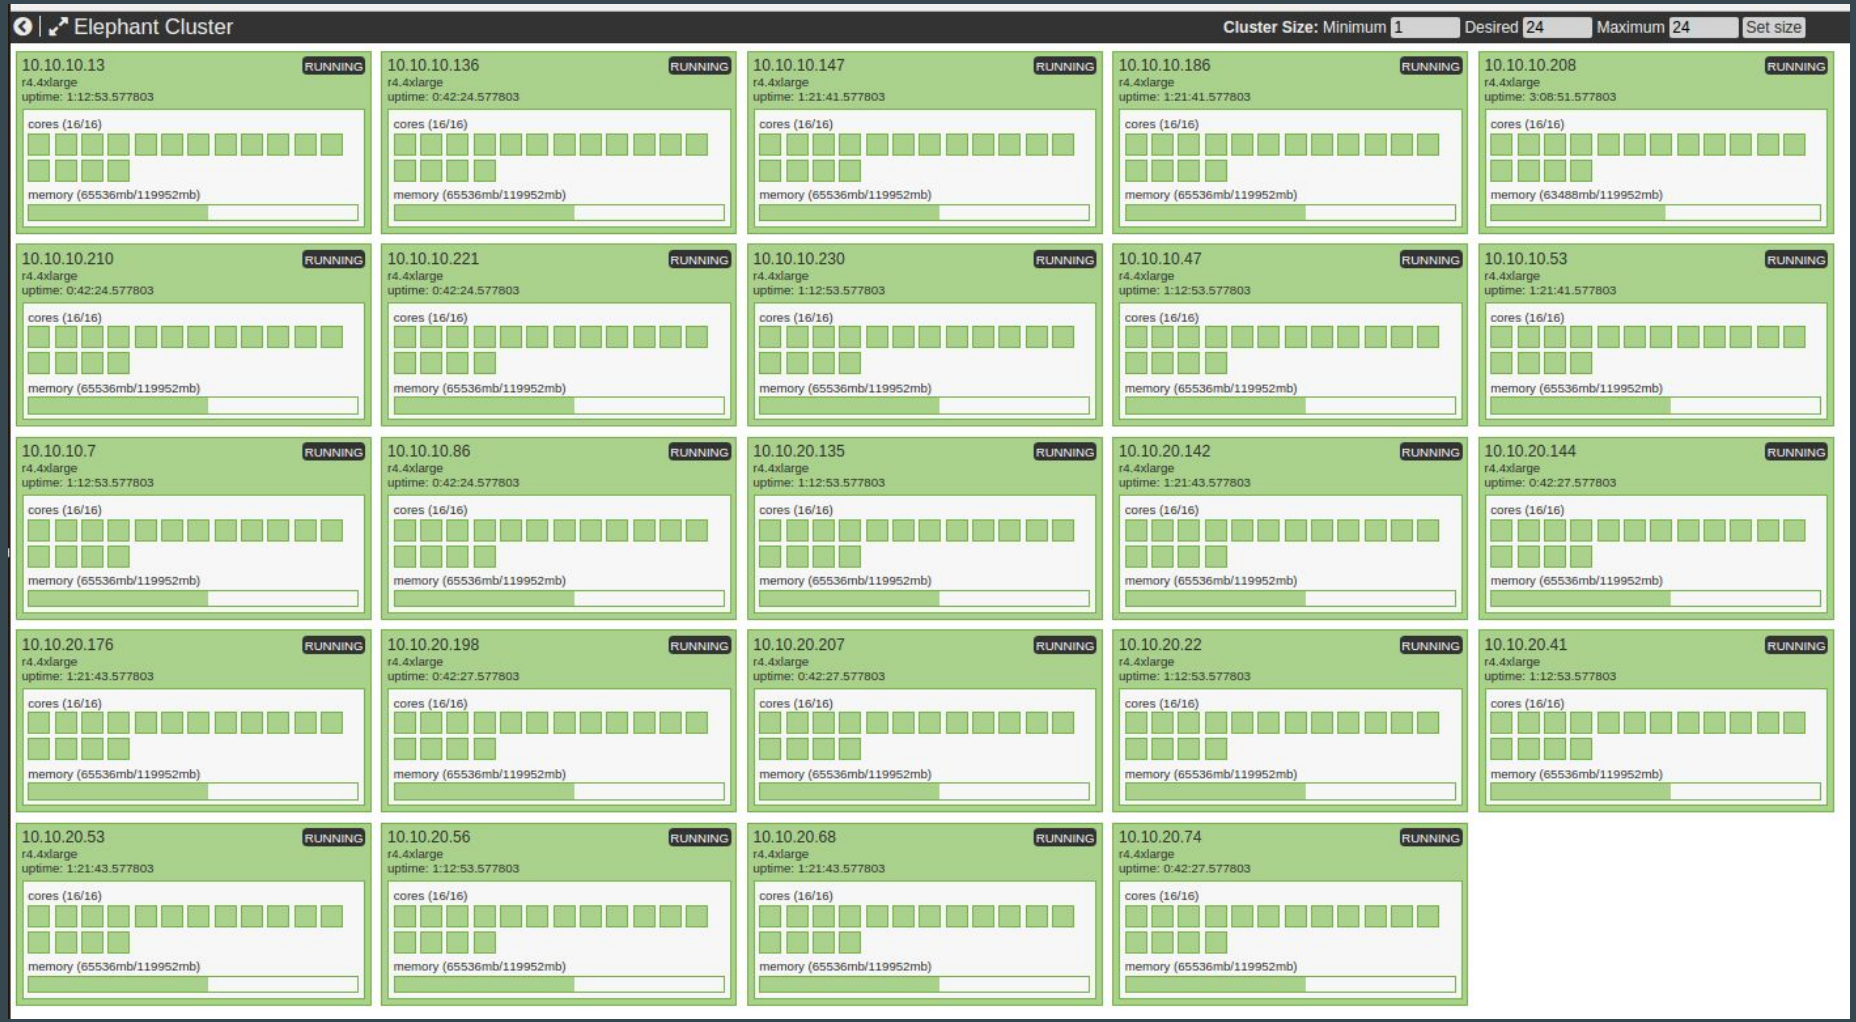
Elephant Cluster
Cluster Size: Minimum 1 Desired 24 Maximum 24 Set size
10.10.10.13 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.136 RUNNING
r4.4xlarge
uptime: 0:42:24.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.147 RUNNING
r4.4xlarge
uptime: 1:21:41.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.186 RUNNING
r4.4xlarge
uptime: 1:21:41.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.208 RUNNING
r4.4xlarge
uptime: 3:08:51.577803
cores (16/16)
memory (63488mb/119952mb)
10.10.10.210 RUNNING
r4.4xlarge
uptime: 0:42:24.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.221 RUNNING
r4.4xlarge
uptime: 0:42:24.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.230 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.47 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.53 RUNNING
r4.4xlarge
uptime: 1:21:41.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.7 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.10.86 RUNNING
r4.4xlarge
uptime: 0:42:24.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.135 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.142 RUNNING
r4.4xlarge
uptime: 1:21:43.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.144 RUNNING
r4.4xlarge
uptime: 0:42:27.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.176 RUNNING
r4.4xlarge
uptime: 1:21:43.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.198 RUNNING
r4.4xlarge
uptime: 0:42:27.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.207 RUNNING
r4.4xlarge
uptime: 0:42:27.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.22 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.41 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.53 RUNNING
r4.4xlarge
uptime: 1:21:43.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.56 RUNNING
r4.4xlarge
uptime: 1:12:53.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.68 RUNNING
r4.4xlarge
uptime: 1:21:43.577803
cores (16/16)
memory (65536mb/119952mb)
10.10.20.74 RUNNING
r4.4xlarge
uptime: 0:42:27.577803
cores (16/16)
memory (65536mb/119952mb)

# Profile the code

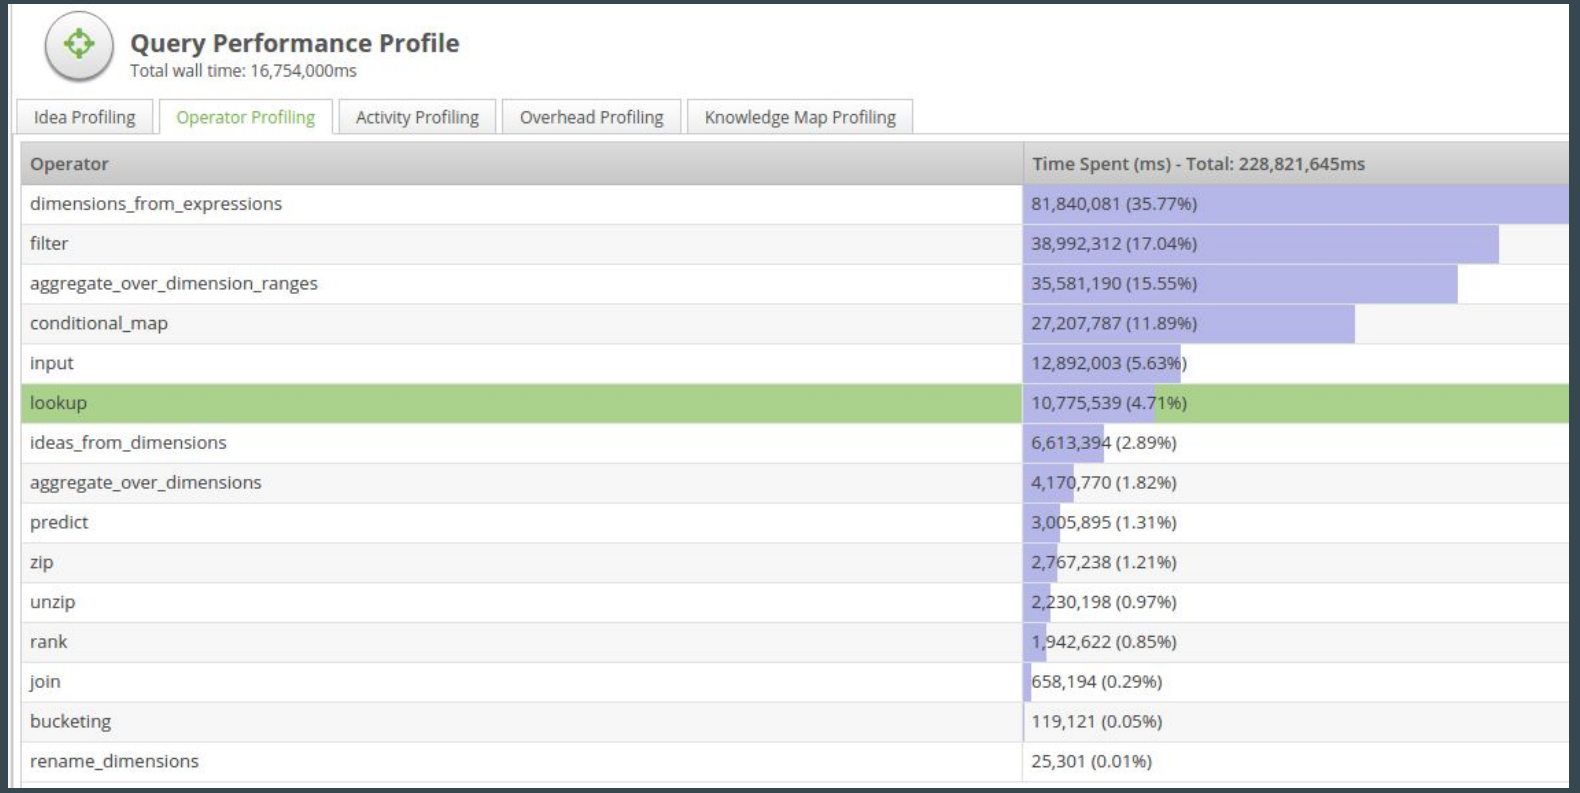

**Query Performance Profile**
Total wall time: 16,754,000ms

Idea Profiling | Operator Profiling | Activity Profiling | Overhead Profiling | Knowledge Map Profiling

| Operator | Time Spent (ms) - Total: 228,821,645ms |
|---|---|
| dimensions_from_expressions | 81,840,081 (35.77%) |
| filter | 38,992,312 (17.04%) |
| aggregate_over_dimension_ranges | 35,581,190 (15.55%) |
| conditional_map | 27,207,787 (11.89%) |
| input | 12,892,003 (5.63%) |
| lookup | 10,775,539 (4.71%) |
| ideas_from_dimensions | 6,613,394 (2.89%) |
| aggregate_over_dimensions | 4,170,770 (1.82%) |
| predict | 3,005,895 (1.31%) |
| zip | 2,767,238 (1.21%) |
| unzip | 2,230,198 (0.97%) |
| rank | 1,942,622 (0.85%) |
| join | 658,194 (0.29%) |
| bucketing | 119,121 (0.05%) |
| rename_dimensions | 25,301 (0.01%) |

# Insights

- Monitoring can makes your life easier
- Compute coming up faster saves $$$ (docker vs puppet): image download vs package install
- Match cluster size to resource requirements
- Having the right autoscaling conf matters

# Autoscaling issues

- Don’t flip on and off
  - Two Causes for us:
    - new job => scale up => complete => scale down => new job (within an instance hour)
    - Scale up too high too quickly => low resource utilization => scale down => (repeat)
- Scale up on load metrics (high cpu), gate scaledown on # messages in alive queue

```
while True:
    try:
        conn = boto.sqs.connect_to_region('us-east-1')
        q = conn.get_queue(heartbeat_queue_name)
        m = Message()
        m.set_body("I am alive")
        q.write(m)
        logger.info("Sending heartbeat message")
        time.sleep(60)
    except Exception: # pylint: disable=broad-except
        logger.exception("Unknown error occurred in heartbeat thread")
        time.sleep(5)
```

# Some recent spot autoscaling usage

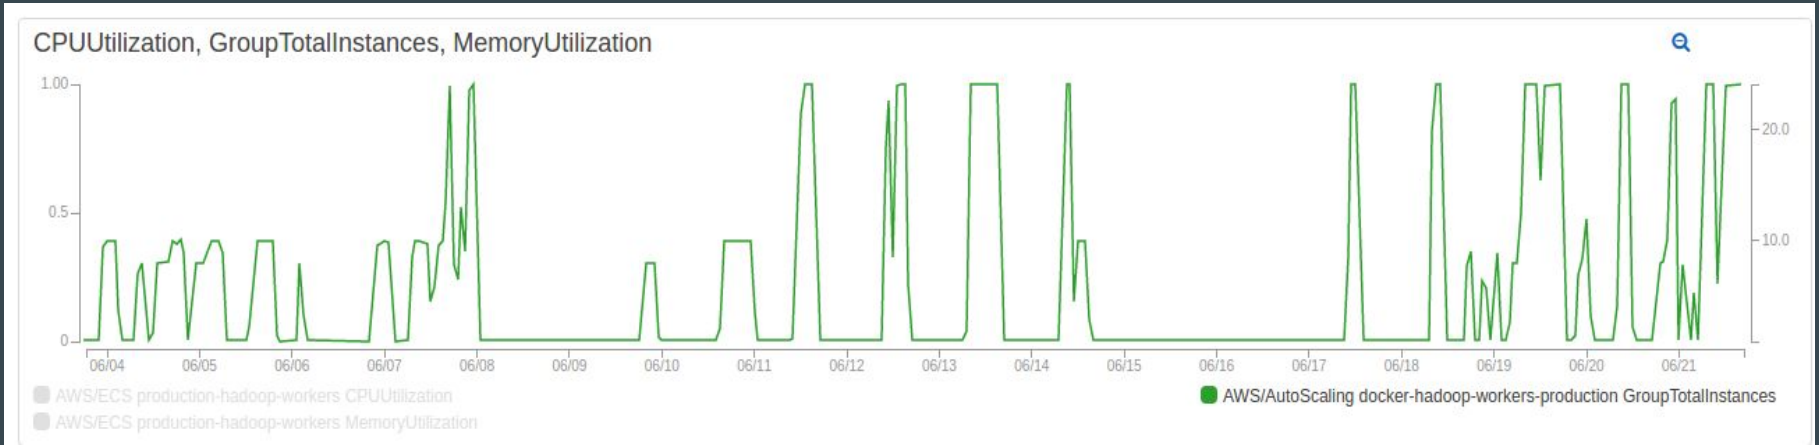

# Questions