

Architecting Failovers

Peter Juritz
peter@nimbula.com

nimbula

- IaaS software for private clouds
- APIs for provisioning and managing infrastructure at scale

What is a failover?

According to wikipedia:

“In computing, failover is automatic switching to a redundant or standby computer server, system, hardware component or network upon the failure or abnormal termination of the previously active application, server, system, hardware component, or network.”

Put more simply

If something breaks, try to fix it automatically.



Who needs failovers?

- Failures WILL happen
- Highly available services
- Operation critical services
- Environments which expect failure
 - At larger scale, expect more failures
- Remember, this is not your main goal

**Engineering
failovers can lead to
Unforeseen
Consequences**

What can cause failures?

- Very difficult to predict, but:
- Bugs in code, both
 - Code you have written
 - Code you rely on

What can cause failures?


- Hardware failures
 - Network
 - Disk
 - Power Outages

What can cause failures?

- The Human Element

Some examples from experience

- The linux leap second bug



Some examples from experience

- Users like to reboot everything



Different cases of failure

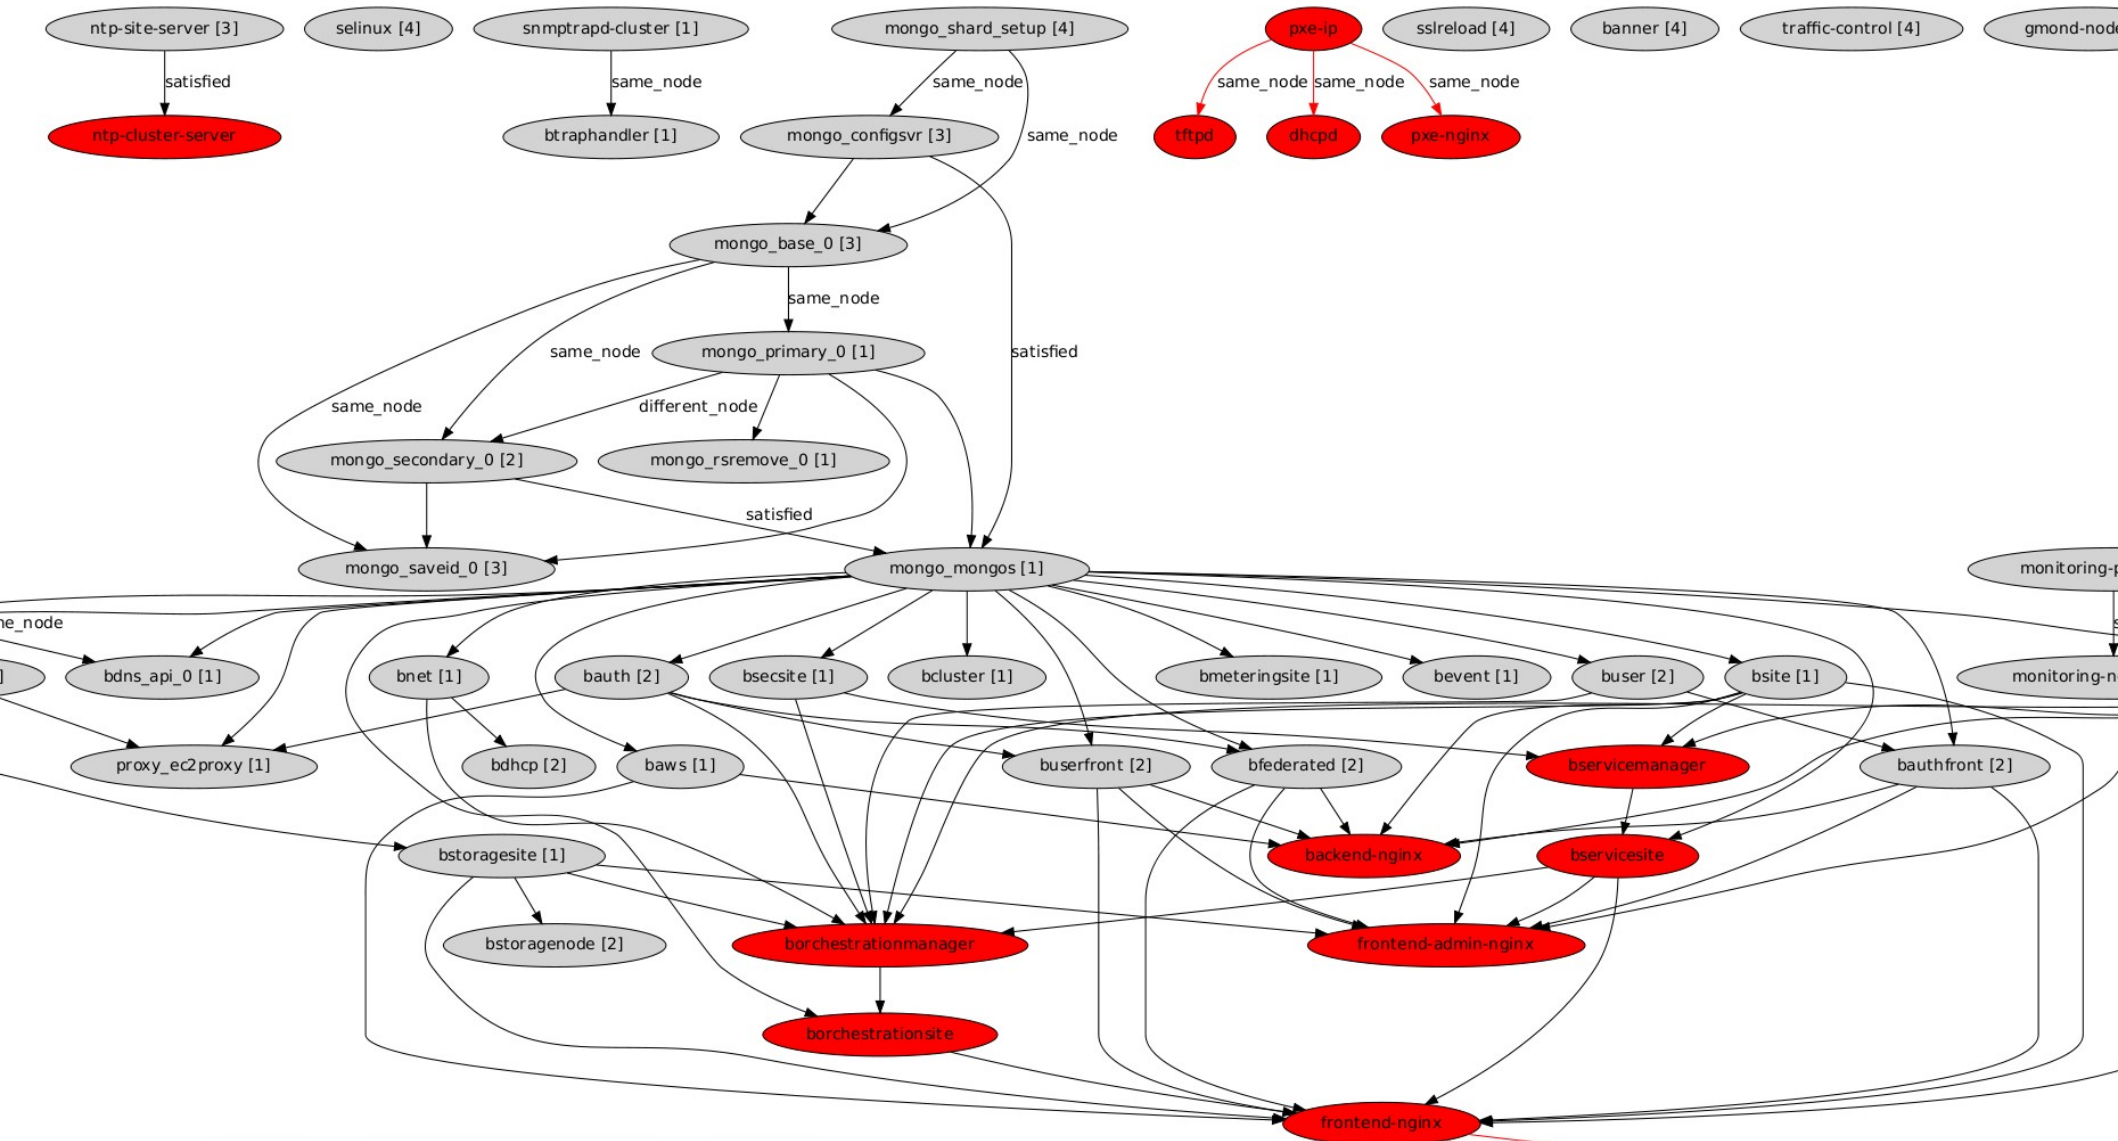
- Hardware dies
- Processes die
- Processes spin/ leak memory
- Data problems (be vary wary)
- Unrecoveravle failures

Detecting a fault

- Not always easy
 - Watch PIDS?
 - Heartbeats? Pings? Idle connections?
 - Health checks?
 - Distributed locks?
 - Very good exception handling/ logging
- This may be the hardest part

Some things we've done

bIC (Infrastructure Controller)



DNS failover

database.internal. → 10.0.0.23

BANG

database.internal. → 10.0.0.42

Apache Zookeeper

- Distributed Coordination system




Simulate Failures

- This can be hard
- You won't catch all the cases

But try! Better now than in
production

Replicated Databases

- Mongo, Riak, Cassandra...
- We happen to use



mongoDB

Some lessons we've
learnt

Some things you can not
recover from – don't waste
time trying to

Simpler service architecture
=
Easier recovery

Thinking about failure forces
you to examine architectural
weakness

Focus your effort on the
most likely problem areas